

Helpfulness (mean reward), higher is better

† rubric-judge cross-check, not comparable with the hand labels

- context arms (B0, B1)
- filter arm and its parent (B2, B3)
- safety-mixed arms (T, T_{ctrl})

better

B0, OR 11.7%†

T, OR 0.0%

T_{ctrl}, OR 23.3%†

B2 and B3 overlap, OR 0.0%

B1, OR 18.3%†

Attack-success rate, primary endpoint (%), lower is better

